

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

c:\Users\bkhan\AppData\Local\Programs\Python\Python39\lib\site-packages\scipy\\_\_init\_\_.py:155: UserWarning: A NumPy version >=1.18.5 and <1.25.0 is required for this version of SciPy (detected version 1.26.1  
warnings.warn(f"A NumPy version >={np\_minversion} and <{np\_maxversion}")

```
In [3]: df = pd.read_csv("Employees Promotion.csv") # loading employee promotion.csv data from day09-project s3 bucket
df.head()
```

```
Out[3]:
```

	EmployeeID	Department	Region_Employment	Education Level	Gender	Recruitment Channel	NO_Trainings_LstYear	Age	previous_year_rating
0	65438	Sales & Marketing	7	Master's & above	f	sourcing	1	35.0	5.0
1	65141	Operations	22	Bachelor's	m	other	1	30.0	5.0
2	7513	Sales & Marketing	19	Bachelor's	m	sourcing	1	34.0	3.0
3	2542	Sales & Marketing	23	Bachelor's	m	other	2	39.0	1.0
4	48945	Technology	26	Bachelor's	m	other	1	45.0	3.0

## Analysing Data

```
In [4]: df.info() # checking data types, dimensions of df
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 54808 entries, 0 to 54807
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   EmployeeID            54808 non-null  int64
 1   Department            54808 non-null  object
 2   Region_Employment     54808 non-null  int64
 3   Education Level       52399 non-null  object
 4   Gender                54808 non-null  object
 5   Recruitment Channel   44404 non-null  object
 6   NO_Trainings_LstYear  54808 non-null  int64
 7   Age                   54268 non-null  float64
 8   previous_year_rating  50684 non-null  float64
 9   Service Length        54808 non-null  int64
10   Awards                54808 non-null  object
11   Avg_Training_Score    52248 non-null  float64
12   Is Promoted           54808 non-null  object
dtypes: float64(3), int64(4), object(6)
memory usage: 5.4+ MB
```

```
In [5]: df.describe() # stats of the df
```

Out[5]:

	EmployeeID	Region_Employment	NO_Trainings_LstYear	Age	previous_year_rating	Service Length	Avg_Training_Score
<b>count</b>	54808.000000	54808.000000	54808.000000	54268.000000	50684.000000	54808.000000	52248.000000
<b>mean</b>	39195.830627	14.195045	1.253011	34.586644	3.329256	5.865512	63.712238
<b>std</b>	22586.581449	10.086273	0.609264	8.114136	1.259993	4.265094	13.521910
<b>min</b>	1.000000	1.000000	1.000000	0.000000	1.000000	1.000000	39.000000
<b>25%</b>	19669.750000	4.000000	1.000000	29.000000	3.000000	3.000000	51.000000
<b>50%</b>	39225.500000	13.000000	1.000000	33.000000	3.000000	5.000000	60.000000
<b>75%</b>	58730.500000	22.000000	1.000000	39.000000	4.000000	7.000000	77.000000
<b>max</b>	78298.000000	34.000000	10.000000	60.000000	5.000000	37.000000	99.000000

In [6]: `df.shape` *# dimesnsions of the df*

Out[6]: (54808, 13)

In [7]: `df.duplicated().sum()` *# checking any duplicated values in the df*

Out[7]: 0

In [8]: `df.isnull().sum()` *# checking total null values in each column*

```
Out[8]: EmployeeID      0
      Department      0
      Region_Employment  0
      Education Level    2409
      Gender           0
      Recruitment Channel 10404
      NO_Trainings_LstYear  0
      Age             540
      previous_year_rating 4124
      Service Length    0
      Awards           0
      Avg_Training_Score 2560
      Is Promoted       0
      dtype: int64
```

In [9]: `df['Is Promoted'].value_counts()` *# target column value counts*

```
Out[9]: NO      50140
      YES      4668
      Name: Is Promoted, dtype: int64
```

## Cleaning and Engineering data

In [10]: `df['Is Promoted'] = df['Is Promoted'].map({'YES': 1, 'NO': 0})` *# mapping the target column values i.e. Yes and No to 1 and 0*In [11]: `df['Education Level'].unique()` *# unique values of Education Level column*

```
Out[11]: array(["Master's & above", "Bachelor's", nan, 'Below Secondary'],
      dtype=object)
```

In [12]: `df['Education Level'].value_counts()` *# Value counts of each values in the Education Level column.*

```
Out[12]: Bachelor's      36669
      Master's & above  14925
      Below Secondary    805
      Name: Education Level, dtype: int64
```

In [13]: `df['Education Level'].isnull().sum()` *# total null values in Education Level column*

Out[13]: 2409

```
In [14]: df['Education Level'].fillna(method='ffill', inplace=True) # using forward fill to fill null values
```

```
In [15]: df['Education Level'].value_counts() # now checking the count
```

```
Out[15]: Bachelor's      38390
Master's & above    15579
Below Secondary      839
Name: Education Level, dtype: int64
```

```
In [16]: df['Education Level'].isnull().sum() # checking null values
```

```
Out[16]: 0
```

```
In [17]: df.Age.unique() # unique values of age column
```

```
Out[17]: array([35., 30., 34., 39., 45., 31., 33., 28., 32., 49., 37., 38., 41.,
        27., 29., 26., 24., 57., 40., 42., 23., 59., nan, 50., 56., 20.,
        25., 47., 36., 46., 44., 60., 0., 43., 22., 54., 58., 48., 53.,
        55., 51., 52., 21.])
```

```
In [18]: df.Age.value_counts() # value counts of age column
```

```
Out[18]: 30.0    3597
32.0    3479
31.0    3474
29.0    3352
33.0    3157
28.0    3103
34.0    3017
27.0    2784
35.0    2673
36.0    2482
37.0    2134
26.0    2019
38.0    1894
39.0    1660
40.0    1634
25.0    1283
41.0    1265
42.0    1133
43.0     978
44.0     834
24.0     832
45.0     749
46.0     686
48.0     554
47.0     550
50.0     516
49.0     435
23.0     419
51.0     386
53.0     354
52.0     345
0.0      340
54.0     306
55.0     285
56.0     262
57.0     234
22.0     228
60.0     214
58.0     211
59.0     202
20.0     110
21.0      98
Name: Age, dtype: int64
```

```
In [19]: df.Age.fillna(df['Age'].median()) # filling the null values of age column using median of age column
```

```
Out[19]: 0      35.0
         1      30.0
         2      34.0
         3      39.0
         4      45.0
         ...
        54803    48.0
        54804    37.0
        54805     0.0
        54806    29.0
        54807    27.0
        Name: Age, Length: 54808, dtype: float64
```

```
In [20]: df.previous_year_rating.unique()    # unique values of previous_year_rating column
```

```
Out[20]: array([ 5.,  3.,  1.,  4., nan,  2.])
```

```
In [21]: df.previous_year_rating.value_counts()    # value counts of previous_year_rating col
```

```
Out[21]: 3.0    18618
         5.0    11741
         4.0     9877
         1.0     6223
         2.0     4225
        Name: previous_year_rating, dtype: int64
```

```
In [22]: df.previous_year_rating.fillna(df.previous_year_rating.median())    # filling previous_year_rating col using its
```

```
Out[22]: 0      5.0
         1      5.0
         2      3.0
         3      1.0
         4      3.0
         ...
        54803    3.0
        54804    2.0
        54805    5.0
        54806    1.0
        54807    1.0
        Name: previous_year_rating, Length: 54808, dtype: float64
```

```
In [23]: df.Avg_Training_Score.unique()    # unique values of Avg_Training_Score col
```

```
Out[23]: array([49., 60., 50., 73., 85., 59., 63., 83., 54., 77., 80., 84., 51.,
                46., 75., 57., 70., 68., 79., 44., 72., nan, 48., 58., 87., 47.,
                52., 88., 71., 65., 62., 53., 78., 91., 82., 69., 55., 74., 86.,
                90., 92., 67., 89., 56., 76., 81., 64., 39., 94., 93., 66., 95.,
                42., 96., 40., 99., 43., 97., 41., 98.])
```

```
In [24]: df.Avg_Training_Score.value_counts()    # value counts of Avg_Training_Score col
```

```
Out[24]: 50.0    2716
         49.0    2681
         48.0    2437
         51.0    2347
         60.0    2155
         59.0    2064
         58.0    1898
         52.0    1856
         47.0    1746
         62.0    1450
         82.0    1447
         57.0    1437
         81.0    1357
         53.0    1324
         80.0    1206
         83.0    1198
         84.0    1168
         79.0    1160
         46.0    1136
         85.0    1072
         56.0    1070
         70.0    1055
         63.0    1021
         69.0    1018
         54.0     997
         68.0     935
         78.0     933
         86.0     912
         71.0     898
         55.0     872
         67.0     728
         72.0     725
         64.0     708
         77.0     697
         87.0     655
         65.0     599
         66.0     580
         73.0     523
         76.0     516
         88.0     444
         74.0     433
         75.0     403
         44.0     335
         89.0     301
         90.0     185
         43.0     176
         91.0     117
         92.0      99
         93.0      84
         94.0      65
         42.0      62
         97.0      49
         96.0      48
         95.0      45
         98.0      37
         99.0      35
         41.0      26
         40.0       5
         39.0       2
Name: Avg_Training_Score, dtype: int64
```

```
In [25]: df.Avg_Training_Score.fillna(df.Avg_Training_Score.median()) # filling null values using Avg_Training_Score r
```

```
Out[25]: 0      49.0
         1      60.0
         2      50.0
         3      50.0
         4      73.0
         ...
        54803    78.0
        54804    56.0
        54805    79.0
        54806    60.0
        54807    49.0
        Name: Avg_Training_Score, Length: 54808, dtype: float64
```

```
In [26]: df = df.drop(['EmployeeID', 'Gender', 'Service Length', 'Recruitment Channel'], axis=1) # dropping these four
```

```
In [27]: df.head()
```

```
Out[27]:
```

	Department	Region_Employment	Education Level	NO_Trainings_LstYear	Age	previous_year_rating	Awards	Avg_Training_Score	Prom
0	Sales & Marketing	7	Master's & above	1	35.0	5.0	NO	49.0	
1	Operations	22	Bachelor's	1	30.0	5.0	NO	60.0	
2	Sales & Marketing	19	Bachelor's	1	34.0	3.0	NO	50.0	
3	Sales & Marketing	23	Bachelor's	2	39.0	1.0	NO	50.0	
4	Technology	26	Bachelor's	1	45.0	3.0	NO	73.0	

```
In [28]: df.duplicated().sum() # checking total duolicate values
```

```
Out[28]: 5013
```

```
In [29]: df.drop_duplicates(keep='first', inplace=True) # dropping duplicates whuile keeping its first value
```

```
In [30]: df.duplicated().sum()
```


```
Out[30]: 0
```

```
In [31]: df['Is Promoted'].value_counts()
```

```
Out[31]: 0    45175
         1    4620
        Name: Is Promoted, dtype: int64
```

```
In [32]: sns.countplot(x=df['Is Promoted'], data=df) # countplot of target variable
```

```
Out[32]: <AxesSubplot:xlabel='Is Promoted', ylabel='count'>
```



```
In [33]: df = pd.get_dummies(df, columns=['Department', 'Education Level', 'Awards']) # getting dummies data of these
```

```
In [34]: df.head()
```

```
Out[34]:
```

	Region_Employment	NO_Trainings_LstYear	Age	previous_year_rating	Avg_Training_Score	Is Promoted	Department_-	Department_
0	7	1	35.0	5.0	49.0	0	0	
1	22	1	30.0	5.0	60.0	0	0	
2	19	1	34.0	3.0	50.0	0	0	
3	23	2	39.0	1.0	50.0	0	0	
4	26	1	45.0	3.0	73.0	0	0	

5 rows × 21 columns

```
In [34]: import os
import boto3
import re
import sagemaker

role = sagemaker.get_execution_role() # execution role
region = boto3.Session().region_name # region name
smclient = boto3.Session().client("sagemaker") # setting sagemaker session
bucket = "day09-project1"
prefix = ("sagemaker/Employess-Promotion") # place to upload training files within the bucket
print (region)
```

```
sagemaker.config INFO - Not applying SDK defaults from location: /etc/xdg/sagemaker/config.yaml
sagemaker.config INFO - Not applying SDK defaults from location: /home/ec2-user/.config/sagemaker/config.yaml
sagemaker.config INFO - Not applying SDK defaults from location: /etc/xdg/sagemaker/config.yaml
sagemaker.config INFO - Not applying SDK defaults from location: /home/ec2-user/.config/sagemaker/config.yaml
us-east-1
```

## Training

```
In [38]: # splitting df into train, validation and test data in 70%, 20% and 10% respectively.
train_data, validation_data, test_data = np.split(df.sample(frac=1, random_state=1729),[int(0.7 * len(df)), int(
pd.concat([train_data["Is Promoted"], train_data.drop("Is Promoted", axis=1)], axis=1).to_csv("train.csv", inde
```

```
pd.concat([validation_data["Is Promoted"], validation_data.drop("Is Promoted", axis=1)], axis=1).to_csv("validation.csv", index=False)
pd.concat([test_data["Is Promoted"], test_data.drop("Is Promoted", axis=1)], axis=1).to_csv("test.csv", index=False)
```

```
In [36]: # storing test and train data
train_y = train_data["Is Promoted"]
train_X = train_data.drop("Is Promoted", axis=1)
test_y = test_data["Is Promoted"]
test_X = test_data.drop("Is Promoted", axis=1)
```

```
In [37]: # uploading train and validation in s3 bucket
boto3.Session().resource("s3").Bucket(bucket).Object(
    os.path.join(prefix, "train/train.csv")
).upload_file("train.csv")
boto3.Session().resource("s3").Bucket(bucket).Object(
    os.path.join(prefix, "validation/validation.csv")
).upload_file("validation.csv")
```

## HP Tuning

ParameterRanges: Specifies the search space for hyperparameters.

- CategoricalParameterRanges: Categorical hyperparameters are those whose values come from a discrete set of categories or options.(none in this example).
- ContinuousParameterRanges: Continuous hyperparameters are those that can take any real value within a specified range:
- eta: Learning rate with a range between 0 and 1.
- min\_child\_weight: Minimum sum of instance weight (hessian) needed in a child.
- alpha: L1 regularization term with a range between 0 and 2.
- IntegerParameterRanges: For integer-valued hyperparameters.
- max\_depth: Maximum depth of a tree with a range between 1 and 10.

ResourceLimits: Specifies resource limits for the tuning job.

- MaxNumberOfTrainingJobs: Maximum number of training jobs to run during the tuning job.
- MaxParallelTrainingJobs: Maximum number of training jobs to run in parallel.

Strategy: Specifies the tuning strategy. In this case, it's set to "Bayesian".

HyperParameterTuningJobObjective: Specifies the objective metric to optimize and whether to minimize or maximize it.

MetricName: Name of the metric to optimize, set to "validation:accuracy" (The ratio of correctly predicted instances to the total instances). Type: Set to "Maximize" since we want to maximize the accuracy.

```
In [38]: from time import gmtime, strftime, sleep

tuning_job_name = "xgboost-tuningjob-" + strftime("%d-%H-%M-%S", gmtime()) # generating unique hyperparameter
print(tuning_job_name)

# configuration for the hyperparameter tuning job
tuning_job_config = {
    "ParameterRanges": {
        "CategoricalParameterRanges": [],
        "ContinuousParameterRanges": [
            {
                "MaxValue": "1",
                "MinValue": "0",
                "Name": "eta", # Learning rate parameter
            },
            {
                "MaxValue": "10",
                "MinValue": "1",
                "Name": "min_child_weight", # Minimum sum of instance weight in a child
            },
            {
                "MaxValue": "2",
                "MinValue": "0",
```



```

        "Name": "alpha",    # L1 regularization term
    },
],
"IntegerParameterRanges": [
    {
        "MaxValue": "10",
        "MinValue": "1",
        "Name": "max_depth",    # Maximum depth of a tree
    }
],
},
"ResourceLimits": {"MaxNumberOfTrainingJobs": 10, "MaxParallelTrainingJobs": 3},
"Strategy": "Bayesian",
"HyperParameterTuningJobObjective": {"MetricName": "validation:accuracy", "Type": "Maximize"},
}

```

xgboost-tuningjob-26-21-56-00

The following block includes details about the algorithm, input data configuration, output data configuration, resource configuration, role ARN, hyperparameters, and stopping conditions.

- The AlgorithmSpecification section specifies the training image and input mode.
- The InputDataConfig section specifies details about the training and validation datasets, including S3 paths.
- The OutputDataConfig section specifies the S3 path where the output model artifacts will be stored.
- The ResourceConfig section specifies the instance count, instance type, and volume size.
- The StaticHyperParameters section specifies static hyperparameters for the XGBoost algorithm.
- The StoppingCondition section specifies the maximum runtime for the training job.

Static hyperparameters:

- `eval_metric`: Specifies the evaluation metric to be used for model performance assessment during training. In this case, the chosen evaluation metric is Area Under the Curve (AUC) for the Receiver Operating Characteristic (ROC) curve. Since it is a common metric for binary classification problems and represents the area under the ROC curve, which measures the model's ability to distinguish between positive and negative examples.
- `num_round`: Specifies the number of boosting rounds (iterations) for the XGBoost algorithm. This parameter determines the number of boosting rounds (iterations) during the training process. A higher number of rounds can potentially lead to a more accurate model, but it's important to monitor for overfitting.
- `objective`: Specifies the learning task and corresponding objective function. This parameter indicates that the training task is binary classification, and the objective function to be optimized is the logistic loss. Since the logistic loss is suitable for binary classification problems, and optimizing it leads to the logistic regression model.
- `rate_drop`: Specifies the rate at which the algorithm drops trees during the dropout process. Dropout is a regularization technique used to prevent overfitting. This parameter controls the fraction of previously dropped trees to be dropped during a training iteration. A value of 0.3 means that 30% of the trees are dropped during each iteration.
- `tweedie_variance_power`: Specifies the power for the variance function in the Tweedie distribution. This parameter is relevant when the objective function is set to Tweedie, which is used for modeling non-negative continuous target variables with a variance function. The specified value of 1.4 determines the power parameter for the Tweedie distribution's variance function.

```

In [39]: from sagemaker.image_uris import retrieve

training_image = retrieve(framework="xgboost", region=region, version="1.5-1")    # Retrieving the URI of the p

# S3 input paths for training and validation data
s3_input_train = "s3://{}/{}/train".format(bucket, prefix)
s3_input_validation = "s3://{}/{}/validation/".format(bucket, prefix)

# configuration for the SageMaker training job
training_job_definition = {
    "AlgorithmSpecification": {"TrainingImage": training_image, "TrainingInputMode": "File"},    # algorithm and
    "InputDataConfig": [    # input data configuration for training and validation
        {
            "ChannelName": "train",    # Channel name for training data
            "CompressionType": "None",
            "ContentType": "csv",

```

```

        "DataSource": {
            "S3DataSource": {
                "S3DataDistributionType": "FullyReplicated",
                "S3DataType": "S3Prefix",
                "S3Uri": s3_input_train,    # S3 path for training data
            },
        },
    },
    {
        "ChannelName": "validation",    # Channel name for validation data
        "CompressionType": "None",
        "ContentType": "csv",
        "DataSource": {
            "S3DataSource": {
                "S3DataDistributionType": "FullyReplicated",
                "S3DataType": "S3Prefix",
                "S3Uri": s3_input_validation,    # S3 path for validation data
            },
        },
    },
],
"OutputDataConfig": {"S3OutputPath": "s3:///{}/output".format(bucket, prefix)},    # output data configuration
"ResourceConfig": {"InstanceCount": 1, "InstanceType": "ml.m4.xlarge", "VolumeSizeInGB": 10},    # resource configuration
"RoleArn": role,    # IAM role
"StaticHyperParameters": {    # static hyperparameters for the XGBoost algorithm
    "eval_metric": "auc",
    "num_round": "100",
    "objective": "binary:logistic",
    "rate_drop": "0.3",
    "tweedie_variance_power": "1.4",
},
"StoppingCondition": {"MaxRuntimeInSeconds": 43200},    # stopping conditions for the training job
}

```

```

In [40]: # creating the hyperparameter tuning job
smclient.create_hyper_parameter_tuning_job(
    HyperParameterTuningJobName=tuning_job_name,
    HyperParameterTuningJobConfig=tuning_job_config,
    TrainingJobDefinition=training_job_definition,
)

```

```

Out[40]: {'HyperParameterTuningJobArn': 'arn:aws:sagemaker:us-east-1:585522057818:hyper-parameter-tuning-job/xgboost-tuningjob-26-21-56-00',
  'ResponseMetadata': {'RequestId': '4764dc56-9e64-411d-b17d-851601b8cfa7',
    'HTTPStatusCode': 200,
    'HTTPHeaders': {'x-amzn-requestid': '4764dc56-9e64-411d-b17d-851601b8cfa7',
      'content-type': 'application/x-amz-json-1.1',
      'content-length': '130',
      'date': 'Sun, 26 Nov 2023 21:56:04 GMT'},
    'RetryAttempts': 0}}

```

```

In [41]: # status of hp
smclient.describe_hyper_parameter_tuning_job(HyperParameterTuningJobName=tuning_job_name)["HyperParameterTuningJobStatus"]

```

```

Out[41]: 'InProgress'

```

```

In [42]: # Job name of hp
smclient.describe_hyper_parameter_tuning_job(HyperParameterTuningJobName=tuning_job_name)["HyperParameterTuningJobName"]

```

```

Out[42]: 'xgboost-tuningjob-26-21-56-00'

```

```

In [50]: # run this cell to check current status of hyperparameter tuning job
tuning_job_result = smclient.describe_hyper_parameter_tuning_job(
    HyperParameterTuningJobName=tuning_job_name
)

status = tuning_job_result["HyperParameterTuningJobStatus"]
if status != "Completed":
    print("Reminder: the tuning job has not been completed.")

job_count = tuning_job_result["TrainingJobStatusCounters"]["Completed"]

```

```
print("%d training jobs have completed" % job_count)

objective = tuning_job_result["HyperParameterTuningJobConfig"]["HyperParameterTuningJobObjective"]
is_minimize = objective["Type"] != "Maximize"
objective_name = objective["MetricName"]

10 training jobs have completed
```

```
In [51]: from pprint import pprint
# extracting best training job from the rest
if tuning_job_result.get("BestTrainingJob", None):
    print("Best model found so far:")
    pprint(tuning_job_result["BestTrainingJob"])
else:
    print("No training jobs have reported results yet.")
```

```
Best model found so far:
{'CreationTime': datetime.datetime(2023, 11, 26, 22, 2, 1, tzinfo=tzlocal()),
 'FinalHyperParameterTuningJobObjectiveMetric': {'MetricName': 'validation:accuracy',
 'Value': 0.9371399879455566},
 'ObjectiveStatus': 'Succeeded',
 'TrainingEndTime': datetime.datetime(2023, 11, 26, 22, 2, 52, tzinfo=tzlocal()),
 'TrainingJobArn': 'arn:aws:sagemaker:us-east-1:585522057818:training-job/xgboost-tuningjob-26-21-56-00-010-62a4af05',
 'TrainingJobName': 'xgboost-tuningjob-26-21-56-00-010-62a4af05',
 'TrainingJobStatus': 'Completed',
 'TrainingStartTime': datetime.datetime(2023, 11, 26, 22, 2, 5, tzinfo=tzlocal()),
 'TunedHyperParameters': {'alpha': '1.524352266875594',
 'eta': '0.5504740010557349',
 'max_depth': '6',
 'min_child_weight': '5.804222989525692'}}
```

```
In [52]: # Now training the model with the best hp values
xgboost_job = "xgboost-" + strftime("%Y-%m-%d-%H-%M-%S", gmtime())
print("Job name is:", xgboost_job)

s3_input_train = "s3://{}/{}/train".format(bucket, prefix)
s3_input_validation = "s3://{}/{}/validation/".format(bucket, prefix)

training_job_definition = {
    "RoleArn": role,
    "TrainingJobName": xgboost_job,
    "AlgorithmSpecification": {"TrainingImage": training_image, "TrainingInputMode": "File"},
    "InputDataConfig": [
        {
            "ChannelName": "train",
            "CompressionType": "None",
            "ContentType": "csv",
            "DataSource": {
                "S3DataSource": {
                    "S3DataDistributionType": "FullyReplicated",
                    "S3DataType": "S3Prefix",
                    "S3Uri": s3_input_train,
                }
            },
        },
        {
            "ChannelName": "validation",
            "CompressionType": "None",
            "ContentType": "csv",
            "DataSource": {
                "S3DataSource": {
                    "S3DataDistributionType": "FullyReplicated",
                    "S3DataType": "S3Prefix",
                    "S3Uri": s3_input_validation,
                }
            },
        },
    ],
    "OutputDataConfig": {"S3OutputPath": "s3://{}/{}/output".format(bucket, prefix)},
    "ResourceConfig": {"InstanceCount": 1, "InstanceType": "ml.m4.xlarge", "VolumeSizeInGB": 10},
    "RoleArn": role,
    "HyperParameters": {
```

```

        "eval_metric": "auc",
        "num_round": "100",
        "objective": "binary:logistic",
        "rate_drop": "0.3",
        "tweedie_variance_power": "1.4",
        "alpha": '1.524352266875594',
        'eta': '0.5504740010557349',
        'max_depth': '6',
        'min_child_weight': '5.804222989525692'
    },
    "StoppingCondition": {"MaxRuntimeInSeconds": 43200},
}

```

Job name is: xgboost-2023-11-26-22-05-31

```

In [53]: sm = boto3.client("sagemaker")
sm.create_training_job(**training_job_definition)
# checking the status of training
status = sm.describe_training_job(TrainingJobName=xgboost_job)["TrainingJobStatus"]
print(status)
sm.get_waiter("training_job_completed_or_stopped").wait(TrainingJobName=xgboost_job)
if status == "Failed":
    message = sm.describe_training_job(TrainingJobName=xgboost_job)["FailureReason"]
    print("Training failed with the following error: {}".format(message))
    raise Exception("Training job failed")

```

InProgress

## Hosting

```

In [54]: # creating a model using the specified configuration.
xgboost_hosting_container = {
    "Image": training_image,
    "ModelDataUrl": sm.describe_training_job(TrainingJobName=xgboost_job)["ModelArtifacts"][
        "S3ModelArtifacts"
    ],
}

create_model_response = sm.create_model(
    ModelName=xgboost_job, ExecutionRoleArn=role, PrimaryContainer=xgboost_hosting_container
)

print(create_model_response["ModelArn"])

```

arn:aws:sagemaker:us-east-1:585522057818:model/xgboost-2023-11-26-22-05-31

```

In [55]: # configuring endpoint with unique name based on current time
xgboost_endpoint_config = "xgboost-endpoint-config-" + strftime("%Y-%m-%d-%H-%M-%S", gmtime())
print(xgboost_endpoint_config)
create_endpoint_config_response = sm.create_endpoint_config(
    EndpointConfigName=xgboost_endpoint_config,
    ProductionVariants=[
        {
            "InstanceType": "ml.m4.xlarge",
            "InitialInstanceCount": 1,
            "ModelName": xgboost_job,
            "VariantName": "AllTraffic",
        }
    ],
)

print("Endpoint Config Arn: " + create_endpoint_config_response["EndpointConfigArn"])

```

xgboost-endpoint-config-2023-11-26-22-10-27

Endpoint Config Arn: arn:aws:sagemaker:us-east-1:585522057818:endpoint-config/xgboost-endpoint-config-2023-11-26-22-10-27

```

In [56]: %%time

# configuring, creating, waiting and then checking the status again of an endpoint
xgboost_endpoint = "xgboost-endpoint-" + strftime("%Y%m%d%H%M", gmtime())
print(xgboost_endpoint)
create_endpoint_response = sm.create_endpoint(

```

```

    EndpointName=xgboost_endpoint, EndpointConfigName=xgboost_endpoint_config
)
print(create_endpoint_response["EndpointArn"])

resp = sm.describe_endpoint(EndpointName=xgboost_endpoint)
status = resp["EndpointStatus"]
print("Status: " + status)

sm.get_waiter("endpoint_in_service").wait(EndpointName=xgboost_endpoint)

resp = sm.describe_endpoint(EndpointName=xgboost_endpoint)
status = resp["EndpointStatus"]
print("Arn: " + resp["EndpointArn"])
print("Status: " + status)

if status != "InService":
    raise Exception("Endpoint creation did not succeed")

xgboost-endpoint-202311262210
arn:aws:sagemaker:us-east-1:585522057818:endpoint/xgboost-endpoint-202311262210
Status: Creating
Arn: arn:aws:sagemaker:us-east-1:585522057818:endpoint/xgboost-endpoint-202311262210
Status: InService
CPU times: user 44.2 ms, sys: 3.02 ms, total: 47.2 ms
Wall time: 3min 31s

```

## Inferencing

```

In [57]: def np2csv(arr):
        csv = io.BytesIO() #the function gets an array (Numpy array) and creates an in-memory binary buffer named csv
        np.savetxt(csv, arr, delimiter=",", fmt="%g") # write the array 'arr' to csv object, columns should be separated by commas
        # In the following line:
        # csv.getvalue() retrieves the entire contents of the buffer csv as a byte string.
        # .decode() converts the byte string into a normal Python string by decoding it using the default UTF-8 encoding.
        # .rstrip() removes any trailing whitespace or newlines from the end of the string.
        return csv.getvalue().decode().rstrip()

```

```

In [58]: import io, json
        runtime = boto3.client("runtime.sagemaker")

        payload = np2csv(test_X)

        # invoking endpoint which is used to make a prediction request to the SageMaker endpoint.
        response = runtime.invoke_endpoint(
            EndpointName=xgboost_endpoint, ContentType="text/csv", Body=payload
        )
        result = response["Body"].read().decode()
        test_pred = np.array(result.split(), dtype=float)
        #The following line:
        # extracts the prediction results from the result dictionary, from the "predictions" key.
        # It is a list of dictionaries where each dictionary has a key "score" representing the model's prediction.

```

```

In [59]: print(test_pred)

[0.98641372 0.00207751 0.03878499 ... 0.18620981 0.92541814 0.20467584]

```

```

In [60]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, roc_curve, confusion_matrix
        import matplotlib.pyplot as plt

        # Assuming test_y and test_pred are your true labels and predicted probabilities

        # Convert probabilities to binary predictions (0 or 1)
        test_pred_class = (test_pred > 0.5).astype(int)

        # Calculate and print accuracy
        accuracy = accuracy_score(test_y, test_pred_class)
        print(f'Accuracy: {accuracy:.4f}')

        # Calculate and print precision, recall, and F1 score
        precision = precision_score(test_y, test_pred_class)
        recall = recall_score(test_y, test_pred_class)

```

```

f1 = f1_score(test_y, test_pred_class)

print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print(f'F1 Score: {f1:.4f}')


# Calculate and print ROC AUC
roc_auc = roc_auc_score(test_y, test_pred)
print(f'ROC AUC: {roc_auc:.4f}')

# Plot ROC curve
fpr, tpr, thresholds = roc_curve(test_y, test_pred)
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()

# Print confusion matrix
conf_matrix = confusion_matrix(test_y, test_pred_class)
print('Confusion Matrix:')
print(conf_matrix)

```

Accuracy: 0.8046  
 Precision: 0.2289  
 Recall: 0.4372  
 F1 Score: 0.3005  
 ROC AUC: 0.7152



Confusion Matrix:  
 [[3798 704]  
 [ 269 209]]

```

In [61]: #If the predicted probability (test_pred) is greater than 0.5, it is considered a prediction for the positive class
#The + 0 part is used to convert the resulting boolean values (True for values above 0.5 and False for values below 0.5) to binary
test_pred_class = (test_pred > 0.5) + 0 # converts the predicted probabilities to binary predictions based on the threshold

# We sort the train_y values and select the median (if is 0, then we select 0 otherwise 1). Then we generate the test_y values

```

```
test_pred_baseline = np.repeat(np.median(train_y), len(test_y)) # This line creates a baseline prediction by i

#compare the binary prediction (test_pred_class) with actual outcomes (test_y)
prediction_accuracy = np.mean((test_y == test_pred_class)) * 100 # calculates the accuracy of your model by c
baseline_accuracy = np.mean((test_y == test_pred_baseline)) * 100 # calculates the accuracy of the baseline by

print("Prediction Accuracy:", round(prediction_accuracy, 1), "%")
print("Baseline Accuracy:", round(baseline_accuracy, 1), "%")
```

Prediction Accuracy: 80.5 %

Baseline Accuracy: 90.4 %

In [62]: `sm.delete_endpoint(EndpointName=xgboost_endpoint)`

Out[62]: {'ResponseMetadata': {'RequestId': '2a85b420-c5c6-4fbb-b9b0-23f16d3516a9',  
'HTTPStatusCode': 200,  
'HTTPHeaders': {'x-amzn-requestid': '2a85b420-c5c6-4fbb-b9b0-23f16d3516a9',  
'content-type': 'application/x-amz-json-1.1',  
'content-length': '0',  
'date': 'Sun, 26 Nov 2023 22:14:34 GMT',  
'connection': 'close'},  
'RetryAttempts': 0}}

In [ ]: